

## Exp5: Installation of Hive on Ubuntu

### Aim:

To Download and install Hive, Understanding Startup scripts, Configuration files.

### Procedure:

#### Step 1: Download and extract it

Download the Apache hive and extract it use tar, the commands given below:

```
$ wget https://downloads.apache.org/hive/hive-3.1.2/apache-hive-3.1.2-bin.tar.gz
```

```
$ tar -xvf apache-hive-3.1.2-bin.tar.gz
```

#### Step 2: Place different configuration properties in Apache Hive

In this step, we are going to do two things ○

Placing Hive Home path in bashrc  
file

```
$ nano .bashrc
```

*And append the below lines in it*

```
export HIVE_HOME=/home/hadoop/apache-hive-3.1.2-bin
export PATH=$PATH:$HIVE_HOME/bin
export HADOOP_USER_CLASSPATH_FIRST=true
```

2. Exporting **Hadoop path in Hive-config.sh** (To communicate with the Hadoop ecosystem we are defining Hadoop Home path in hive config field) **Open the**

**hiveconfig.sh as shown in below** *\$ cd apache-hive-3.1.2-bin/bin*

```
$ cp hive-env.sh.template hive-env.sh
```

```
$ nano hive-env.sh
```

*Append the below commands on it* **export**

```
HADOOP_HOME=/home/Hadoop/Hadoop
```

```
export HIVE_CONF_DIR=/home/Hadoop/apache-hive-3.1.2/conf
```

```
# Set HADOOP_HOME to point to a specific hadoop install directory
# HADOOP_HOME=${bin}/../..../hadoop
export HADOOP_HOME=/home/hadoop/hadoop

# Hive Configuration Directory can be controlled by:
# export HIVE_CONF_DIR=
export HIVE_CONF_DIR=/home/hadoop/apache-hive-3.1.2-bin/conf
# Folder containing extra libraries required for hive compilation/execution can be controlled by:
```

#### Step 3: Install mysql

1. Install mysql in Ubuntu by running this command:

```
$ sudo apt update
```

```
$ sudo apt install mysql-server
```

2. *Alter username and password for MySQL by running below commands:*

```
$ sudo mysql
```

Pops command line interface for MySQL and run the below SQL queries to change username and set password

```
mysql> SELECT user, host, plugin FROM mysql.user WHERE user = 'root';
```

```
mysql> ALTER USER 'root'@'localhost' IDENTIFIED WITH 'mysql_native_password' BY  
'your_new_password';
```

```
mysql> FLUSH PRIVILEGES;
```

#### **Step 4: Config hive-site.xml**

Config the hive-site.xml by appending this xml code and change the username and password according to your MySQL.

```
$cd apache-hive-3.1.2-bin/bin
```

```
$cp hive-default.xml.template hive-site.xml
```

```
$nano hive-site.xml
```

*Append these lines into it*

*Replace root as your username of MySQL*

*Replace your\_new\_password as with your password of MySQL*

```
<configuration>
```

```
<property>
```

```
<name>javax.jdo.option.ConnectionURL</name>
```

```
<value>jdbc:mysql://localhost/metastore?createDatabaseIfNotExist=true</value>
```

```
</property>
```

```
<property>
```

```
<name>javax.jdo.option.ConnectionDriverName</name>
```

```
<value>com.mysql.cj.jdbc.Driver</value>
```

```
</property>
```

```
<property>
```

```
<name>javax.jdo.option.ConnectionUserName</name>
```

```
<value>root</value>
```

```
</property>
```

```
<property>
```

```
<name>javax.jdo.option.ConnectionPassword</name>
```

```
<value>your_new_password</value>
```

```
</property>
```

```
<property>
```

```
<name>datanucleus.autoCreateSchema</name>
```

```
<value>true</value>
```

```
</property>
```

```
<property>
```

```
<name>datanucleus.fixedDatastore</name>
```

```
<value>true</value>
```

```
</property>

<property>
<name>datanucleus.autoCreateTables</name>
<value>True</value>
</property>
```

```
</configuration>
```

### Step 5: Setup MySQL java connector:

First, you'll need to download the MySQL Connector/J, which is the JDBC driver for MySQL. You can download it from the below link

[https://drive.google.com/file/d/1QFhB7Kvc7a4LzDRe6GcmZva1yAxKz/view?usp=drive\\_link](https://drive.google.com/file/d/1QFhB7Kvc7a4LzDRe6GcmZva1yAxKz/view?usp=drive_link)

Copy the downloaded MySQL Connector/J JAR file to the Hive library directory. By default, the Hive library directory is usually located at `/path/to/apache-hive-3.1.2/lib/` on Ubuntu. Use the following command to copy the JAR file:

`$sudo cp /path/to/mysql-connector-java-8.0.15.jar /path/to/apache-hive-3.1.2/lib/` Replace `/path/to/` with the actual path to the JAR file.

### Step 6: Initialize the Hive Metastore Schema:

Run the following command to initialize the Hive metastore schema:

`$$HIVE_HOME/bin/schematool -initSchema -dbTypemysql`

### Step 7: Start hive:

You can test Hive by running the Hive shell: Copy code hive You should be able to run Hive queries, and metadata will be stored in your MySQL database. `$hive`

```
hadoop@ubuntu:~/Ex5$ hive
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/hadoop/apache-hive-3.1.3-bin/lib/log4j-slf4j-impl-2.17.1.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/hadoop/hadoop/share/hadoop/common/lib/slf4j-reload4j-1.7.36.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
Hive Session ID = 18ce7f00-0b72-4bd8-a322-9841ea551f4c

Logging initialized using configuration in jar:file:/home/hadoop/apache-hive-3.1.3-bin/lib/hive-common-3.1.3.jar!/hive-log4j2.properties Async: true
Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e . spark, tez) or using Hive 1.X releases.
Hive Session ID = ffb7a4cc-68c4-4355-8a1f-97ec82c24f87
hive> █
```

### Result:

Thus, the Apache Hive installation is completed successfully on Ubuntu.

## Exp5a: Design and test various schema models to optimize data storage and retrieval Using Hive.

### Aim:

To Design and test various schema models to optimize data storage and retrieval Using Hbase.

### Procedure:

#### Step 1: Start Hive

Open a terminal and start Hive by running:

**\$hive**

#### Step 2: Create a Database

Create a new database in Hive: **hive>CREATE DATABASE financials;**

```
hive> DROP DATABASE financials CASCADE;
OK
Time taken: 2.629 seconds
hive> show databases;
OK
default
Time taken: 0.083 seconds, Fetched: 1 row(s)
hive> CREATE DATABASE financials;
OK
Time taken: 0.307 seconds
hive> 
```

#### Step 3: Use the Database:

Switch to the newly created database: **hive>use financials;**

```
hive> use financials;
OK
Time taken: 0.116 seconds
hive> CREATE TABLE finance_table( id INT, name STRING );
OK
Time taken: 1.171 seconds
```

#### Step 4: Create a Table:

Create a simple table in your database:

**hive>CREATE TABLE finance\_table( id INT, name STRING );**

```
hive> use financials;
OK
Time taken: 0.116 seconds
hive> CREATE TABLE finance_table( id INT, name STRING );
OK
Time taken: 1.171 seconds
```

#### Step 5: Load Sample Data:

You can insert sample data into the table:

**hive>INSERT INTO finance\_tableVALUES (1, 'Alice'), (2, 'Bob'), (3, 'Charlie');**

```

hive> INSERT INTO finance_table VALUES (1, 'Alice'), (2, 'Bob'), (3, 'Charlie');
Query ID = hadoop_20240920194248_2da89626-c11d-4c1d-8f6c-b40eda1ba05d
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Job running in-process (local Hadoop)
2024-09-20 19:42:58,605 Stage-1 map = 0%,   reduce = 0%
2024-09-20 19:43:02,334 Stage-1 map = 100%,   reduce = 100%
Ended Job = job_local768369549_0001
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to directory hdfs://localhost:9000/user/hive/warehouse/financials.db/finance_table/.hive-staging_hive_2024-09-20_19-42-4
8_631_1423264935240127676-1/-ext-10000
Loading data to table financials.finance_table
MapReduce Jobs Launched:
Stage-Stage-1:   HDFS Read: 0 HDFS Write: 208 SUCCESS
Total MapReduce CPU Time Spent: 0 msec
OK
Time taken: 15.944 seconds

```

### Step 6: Query Your Data

Use SQL-like queries to retrieve data from your table:

*hive>CREATE VIEW myview AS SELECT name, id FROM finance\_table;*

### Step 7: View the data:

To see the data in the view, you would need to query the view *hive>SELECT\*FROM myview;*

```

hive> CREATE VIEW myview AS SELECT name, id FROM finance_table;
OK
Time taken: 0.558 seconds
hive> SELECT * FROM myview;
OK
Alice      1
Bob         2
Charlie     3
Time taken: 0.428 seconds, Fetched: 3 row(s)

```

### Step 8: Describe a Table:

You can describe the structure of a table using the DESCRIBE command:

*hive>DESCRIBE finance\_table;*

```

hive> DESCRIBE finance_table;
OK
id                int
name              string
Time taken: 0.211 seconds, Fetched: 2 row(s)
hive> ALTER TABLE finance_table ADD COLUMNS (age INT);
OK
Time taken: 0.409 seconds
hive> quit;

```

### Step 9: Alter a Table:

You can alter the table structure by adding a new column: *hive>ALTER TABLE finance\_table ADD COLUMNS (age INT);*

```

hive> DESCRIBE finance_table;
OK
id                int
name              string
Time taken: 0.211 seconds, Fetched: 2 row(s)
hive> ALTER TABLE finance_table ADD COLUMNS (age INT);
OK
Time taken: 0.409 seconds
hive> quit;

```

### Step 10: Quit Hive:

To exit the Hive CLI, simply type: *hive>quit;*

### Result:

Thus, the usage of various commands in Hive has been successfully completed.