

For hadoop.3x version

http://localhost:9870

http://localhost:8088/cluster

Place the mapper file ,reducer file and hadoop dtarming jar file in Documents,Create input foler in hadoop and place the wordcount.txt file on it.

hadoop/Documents\$ give below comments to run

```
hadoop@Ubuntu:~/Documents$ hadoop jar hadoop-streaming-2.7.3.jar -input
/home/hadoop/input/word_count_data.txt -output /home/hadoop/output -mapper mapper.py
reducer reducer.py
```

To check the output folder part-oooo file is created or not

```
hadoop@Ubuntu:~/Documents$ hadoop fs -ls /home/hadoop/output
```

```
hadoop@Ubuntu:~/Documents$ hadoop fs -ls /home/hadoop/output
```

Found 2 items

```
-rw-r--r-- 1 hadoop supergroup      0 2024-08-03 08:59 /home/hadoop/output/_SUCCESS
-rw-r--r-- 1 hadoop supergroup    592 2024-08-03 08:59 /home/hadoop/output/part-00000
```

```
hadoop@Ubuntu:~/Documents$ hdfs dfs -cat /user/hadoop/output/part-00000
```

cat: `/user/hadoop/output/part-00000': No such file or directory

Verify the output

```
hadoop@Ubuntu:~/Documents$ hdfs dfs -cat /home/hadoop/output/part-00000
```

2,000 1

ChatGPT 1

Did 1

Roman 2

Romans 1

Some 1

Sure! 1

This 1

a 3

actually 1

ancient 1 and

3

ash 1 ash,

1

because
1
buildings
1
called 1
concrete 2 concrete, 1
concrete. 1
construction 1 durable
1 for 1 form 2 from 1
gets 1 harbors 1
has 1 have 1
impressive 1
incredibly 1
is 1 know
1 lime, 1
longevity
1 made
1 many 1
mineral
1
mixture
1
modern
1 of
3
outlasted
1 over
2 partly
1 reacts 1
reinforces
1
seawater
1
seawater,
1
showcasing
1
stronger
1
structures?
1
survived
1
techniques
1
that 2
the 5
their 1

time. 1 to
1
tobermorite, 1
used 1
volcanic 2
was 1
which 1
with 1
years, 1 you
1

path of hadoop input file
/home/hadoop/input/word_count_data.txt

path of hadoop output file /home/hadoop/output

Commands on hadoop to check the input and output file

1. List Contents of a Hadoop Directory

To list the contents of a directory in HDFS, use the `-ls` option with the `hadoop fs` command.

List the Input Folder

bash

`hadoop fs -ls /user/hduser/input`

List the Output Folder

bash

`hadoop fs -ls /user/hduser/output`

2. View Detailed Information

The `-ls` command provides a detailed listing of files and directories, including permissions, owner, group, size, and modification date.

3. View File Contents

To view the contents of a file, use the `-cat` option. For example:

View a File in the Input Folder

```
bash
```

```
hadoop fs -cat /user/hduser/input/filename.txt
```

Replace `filename.txt` with the actual name of the file you want to view.

View a File in the Output Folder

If your output folder contains multiple files (e.g., `part-r-00000`), you can view one of the files:

```
bash
```

```
hadoop fs -cat /user/hduser/output/part-r-00000
```

4. Check for Folder Existence

To check if a folder exists in HDFS, you can use the `-test` command with the `-d` option:

Check if Input Folder Exists

```
bash
```

```
hadoop fs -test -d /user/hduser/input && echo "Input folder exists" || echo "Input folder does not exist"
```

Check if Output Folder Exists

```
bash
```

```
hadoop fs -test -d /user/hduser/output && echo "Output folder exists" || echo "Output folder does not exist"
```

5. Additional Commands

Make a Directory:

```
bash
```

```
hadoop fs -mkdir /user/hduser/new_directory
```

Delete a File or Directory:

```
bash
```

```
hadoop fs -rm /user/hduser/input/filename.txt
```

```
hadoop fs -rm -r /user/hduser/output
```

Get File Status:

```
bash
```

```
hadoop fs -stat /user/hduser/input/filename.txt
```

Example Usage

Assuming you have the following setup:

Input Folder: /user/hduser/input

Output Folder: /user/hduser/output

You would list the contents like this:

```
bash
```

```
hadoop fs -ls /user/hduser/input
```

And to check the output:

```
bash
```

```
hadoop fs -ls /user/hduser/output
```

```
bash
```

```
hadoop fs -cat /user/hduser/output/part-r-00000
```

PIG UDF PROGRAM

To check the pig program

```
hadoop@Ubuntu:~/Documents$ nano sample.txt
```

Paste the below content to sample.txt

```
1,John
```

```
2,Jane
```

```
3,Joe
```

```
4,Emma
```

```
hadoop@Ubuntu:~/Documents$ hadoop fs -put sample.txt /home/hadoop/piginput/
```

```
hadoop@Ubuntu:~/Documents$ nano demo_pig.pig
```

paste the below the content to demo_pig.pig

-- Load the data from HDFS

```
data = LOAD '/home/hadoop/piginput/sample.txt' USING PigStorage(',') AS (id:int>
```

-- Dump the data to check if it was loaded correctly

```
DUMP data;
```

```
hadoop@Ubuntu:~/Documents$ pig demo_pig.pig
```

```
2024-08-07 12:13:08,791 [main] INFO
```

```
org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
```

```
(1,John)
```

```
(2,Jane)
```

```
(3,Joe)
```

```
(4,Emma)
```

By using these commands, you can manage and inspect files and directories in your Hadoop setup.

-----up-----

To Run pig basic program and uf program

```
uppercase_udf.py
```

```
- def uppercase(text):    return text.upper()
```

```
if __name__ ==
```

```
"__main__":    import sys
```

```
for line in sys.stdin:
```

```
    line = line.strip()    result
```

```
= uppercase(line)
print(result)
```

Create the udfs folder on hadoop

```
hadoop@Ubuntu:~/Documents$ hadoop fs -mkdir /home/hadoop/udfs
```

put the upppercase_udf.py in to the abv folder hadoop@Ubuntu:~/Documents\$ hdfs

```
dfs -put uppercase_udf.py /home/hadoop/udfs/ hadoop@Ubuntu:~/Documents$
```

```
nano udf_example.pig udf_example.pig
```

-- Register the Python UDF script

```
REGISTER 'hdfs:///home/hadoop/udfs/uppercase_udf.py' USING jython AS udf;
```

-- Load some data data = LOAD 'hdfs:///home/hadoop/sample.txt'

```
AS (text:chararray);
```

-- Use the Python UDF

```
uppercased_data = FOREACH data GENERATE udf.uppercase(text) AS uppercase_text;
```

-- Store the result

```
STORE uppercased_data INTO 'hdfs:///home/hadoop/pig_output_data';
```

place sample.txt file on hadoop hadoop@Ubuntu:~/Documents\$ hadoop

```
fs -put sample.txt /home/hadoop/
```

To Run the pig file

```
hadoop@Ubuntu:~/Documents$ pig -f udf_example.pig
```


finally u get

Success!

Job Stats (time in seconds):

| JobId | Maps | Reduces | MaxMapTime | MinMapTime | AvgMapTime | MedianMapTime | MaxReduceTime | MinReduceTime | AvgReduceTime | MedianReducetime | Alias | Feature | Outputs |
|--|------|---------|------------|------------|------------|---------------|---------------|---------------|---------------|------------------|-------|---------|---------|
| job_local1786848041_0001 | 1 | 0 | n/a | n/a | n/a | n/a | 00 | 0 | 0 | | | | |
| data,uppercased_data MAP_ONLY hdfs:///home/hadoop/pig_output_data, | | | | | | | | | | | | | |

Input(s):

Successfully read 4 records (42778068 bytes) from: "hdfs:///home/hadoop/sample.txt"

Output(s):

Successfully stored 4 records (42777870 bytes) in: "hdfs:///home/hadoop/pig_output_data"

Counters:

Total records written : 4

Total bytes written : 42777870

Spillable Memory Manager spill count : 0

Total bags proactively spilled: 0

Total records proactively spilled: 0

Job DAG:

job_local1786848041_0001

2024-08-07 13:33:04,631 [main] WARN org.apache.hadoop.metrics2.impl.MetricsSystemImpl - JobTracker metrics system already initialized!

2024-08-07 13:33:04,639 [main] WARN org.apache.hadoop.metrics2.impl.MetricsSystemImpl - JobTracker metrics system already initialized!

2024-08-07 13:33:04,644 [main] WARN org.apache.hadoop.metrics2.impl.MetricsSystemImpl - JobTracker metrics system already initialized!

2024-08-07 13:33:04,667 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Success!
To check the output file is created

```
hadoop@Ubuntu:~/Documents$ hdfs dfs -ls /home/hadoop/pig_output_data
```

Found 2 items

If you need to examine the files in the output folder, use:

To view the output

```
hadoop@Ubuntu:~/Documents$ hdfs dfs -cat /home/hadoop/pig_output_data/part-m-00000
```

1,JOHN

2,JANE

3,JOE

4,EMMA

Create json file on bash & save as emp.json

nano emp.json ; Paste the below content on it

```
[
  {"name": "John Doe", "age": 30, "department": "HR", "salary": 50000},
  {"name": "Jane Smith", "age": 25, "department": "IT", "salary": 60000},
  {"name": "Alice Johnson", "age": 35, "department": "Finance", "salary": 70000},
  {"name": "Bob Brown", "age": 28, "department": "Marketing", "salary": 55000},
  {"name": "Charlie Black", "age": 45, "department": "IT", "salary": 80000}
]
```

Check json is readable or any error by giving **install**

jq by **sudo apt-get install jq** **hadoop@Ubuntu:~\$**

jq . emp.json

```
[
  {
    "name": "John Doe",
    "age": 30,
    "department": "HR",
    "salary": 50000
  },
  {
```

```
"name": "Jane Smith",
"age": 25,
"department": "IT",
"salary": 60000
},
{
  "name": "Alice Johnson",
  "age": 35,
  "department": "Finance",
  "salary": 70000
},
{
  "name": "Bob Brown",
  "age": 28,
  "department": "Marketing",
  "salary": 55000
},
{
  "name": "Charlie Black",
  "age": 45,
  "department": "IT",
  "salary": 80000
}
]
```

bash: put the employees.json local directory to *home/hadoop* directory

Example

Suppose the original employees relation has the following data:

| name | age | department | salary |
|---------------|-----|------------|--------|
| John Doe | 30 | HR | 50000 |
| Jane Smith | 25 | IT | 60000 |
| Alice Johnson | 35 | Finance | 70000 |

Bob Brown 28 Marketing 55000
Charlie Black 45IT 80000

After executing:

pig shell: Load the json file by giving following command

```
grunt>-- Load the data employees = LOAD '/home/hadoop/emp.json' USING
```

```
JsonLoader('name:chararray,age:int,department:chararray,salary:float'); grunt>projected
```

```
= FOREACH employees GENERATE name, salary;
```

```
DUMP projected;
```

The projected relation will look like:

| name | salary |
|---------------|---------------|
| John Doe | 50000 |
| Jane Smith | 60000 |
| Alice Johnson | 70000 |
| Bob Brown | 55000 |
| Charlie Black | 80000 |

Assume your employees dataset looks like this:

| name | age | department | salary |
|---------------|-----|------------|--------|
| John Doe | 30 | HR | 50000 |
| Jane Smith | 25 | IT | 60000 |
| Alice Johnson | 35 | Finance | 70000 |
| Bob Brown | 28 | Marketing | 55000 |
| Charlie Black | 45 | IT | 80000 |

1. Aggregation

Aggregate the total salary:

pig

-- Load the data

```
employees = LOAD '/home/hadoop/employees.json' USING
```

```
JsonLoader('name:chararray,age:int,department:chararray,salary:float');
```

-- Aggregate: Calculate the total salary

```
total_salary = FOREACH (GROUP employees ALL) GENERATE SUM(employees.salary) AS  
total_salary;
```

```
DUMP total_salary;
```

Output:

scss

(315000.0)

2. Skip

Skip the first 2 records:

pig

-- Load the data

```
employees = LOAD '/home/hadoop/employees.json' USING  
JsonLoader('name:chararray,age:int,department:chararray,salary:float');
```

-- Skip the first 2 records skipped_employees = LIMIT employees 1000000; -- Use
LIMIT to handle skipping

```
DUMP skipped_employees;
```

Output:

| name | age | department | salary |
|---------------|-----|------------|--------|
| Alice Johnson | 35 | Finance | 70000 |
| Bob Brown | 28 | Marketing | 55000 |
| Charlie Black | 45 | IT | 80000 |

Note: The LIMIT command should be used with an appropriate number, as Pig does not directly support skipping a specific number of records.

3. Limit

Limit the results to the top 3 records:

pig

-- Load the data

```
employees = LOAD '/home/hadoop/employees.json' USING  
JsonLoader('name:chararray,age:int,department:chararray,salary:float');
```

-- Limit: Get the top 3 highest earners

```
top_3_employees = LIMIT employees 3;
```

```
DUMP top_3_employees;
```

Output:

| name | age | department | salary |
|---------------|-----|------------|--------|
| Charlie Black | 45 | IT | 80000 |
| Alice Johnson | 35 | Finance | 70000 |
| Jane Smith | 25 | IT | 60000 |

4. Count

Count the number of employees:

pig

-- Load the data

```
employees = LOAD '/home/hadoop/employees.json' USING  
JsonLoader('name:chararray,age:int,department:chararray,salary:float');
```

-- Count the number of employees

```
employee_count = FOREACH (GROUP employees ALL) GENERATE COUNT(employees) AS  
total_count;
```

```
DUMP employee_count;
```

Output:

```
scss
```

(5)

5. Remove

Remove employees from a specific department, e.g., "IT":

```
pig
```

```
-- Load the data
```

```
employees = LOAD '/home/hadoop/employees.json' USING  
JsonLoader('name:chararray,age:int,department:chararray,salary:float');
```

```
-- Remove employees from the 'IT' department
```

```
filtered_employees = FILTER employees BY department != 'IT';
```

```
DUMP filtered_employees;
```

Output:

| name | age | department | salary |
|---------------|-----|------------|--------|
| John Doe | 30 | HR | 50000 |
| Alice Johnson | 35 | Finance | 70000 |
| Bob Brown | 28 | Marketing | 55000 |

=====

import Json file and do projection, aggregation, limit,count ,skip and remove using python and hdfs

Steps to be followed:

Install pandas and hdfs using pip.

- **Optionally** install pyarrow or hdfs3 if needed based on your specific requirements.
- **Verify** the installation to ensure everything is set up correctly.

Required Packages

pandas:

Purpose: Provides data structures and functions to efficiently manipulate and analyze data.

Installation: Use pip to install pandas.

bash

pip install pandas

hdfs:

Purpose: Provides a Python interface to interact with HDFS.

Installation: Use pip to install hdfs.

bash

pip install hdfs

Additional Considerations

While the script should work with just the above packages, here are some additional considerations:

pyarrow (Optional but useful):

Purpose: If you're working with Apache Arrow or need additional features for handling large datasets or different file formats, pyarrow can be useful.

Installation: Use pip to install pyarrow.

```
bash
```

```
pip install pyarrow
```

hdfs3 (Alternative to hdfs):

Purpose: Another Python library for interacting with HDFS. It's an alternative to the hdfs package and might be preferred in some scenarios.

Installation: Use pip to install hdfs3.

```
bash
```

```
pip install hdfs3
```

Verifying Package Installation

After installing the required packages, you can verify that they are correctly installed and accessible in your Python environment:

```
python
```

```
import pandas as pd from hdfs
```

```
import InsecureClient
```

```
# Check pandas version
```

```
print("Pandas version:", pd.__version__)
```

```
# Test HDFS client connection client =
InsecureClient('http://localhost:9870', user='hadoop')
print("HDFS status:", client.status('/'))
```

If you run this script and see the version of pandas and a status message from HDFS without any errors, the packages are installed correctly.

Create process_data.py file

```
from hdfs import
```

```
InsecureClient import pandas
```

```
as pd import json
```

```
# Connect to HDFS hdfs_client =
```

```
InsecureClient('http://localhost:9870', user='hdfs')
```

```
# Read JSON data from HDFS try: with
```

```
hdfs_client.read('/home/hadoop/emp.json', encoding='utf-8') as reader:
```

```
    json_data = reader.read() # Read the raw data as a
```

```
string    if not json_data.strip(): # Check if data is empty
```

```
raise ValueError("The JSON file is empty.")
```

```
    print(f'Raw JSON Data: {json_data[:1000]}') # Print first 1000 characters for
```

```
debugging    data = json.loads(json_data) # Load the JSON data except
```

```
json.JSONDecodeError as e: print(f'JSON Decode Error: {e}') exit(1) except Exception
```

```
as e:
```

```
    print(f'Error reading or parsing JSON data: {e}')
exit(1)
```

```
# Convert JSON data to DataFrame
```

```
try:
```

```
    df = pd.DataFrame(data)
```

```
except ValueError as e:
```

```
print(f'Error converting JSON data to DataFrame: {e}')
exit(1)

# Projection: Select only 'name' and 'salary' columns
projected_df = df[['name', 'salary']]

# Aggregation: Calculate total salary
total_salary = df['salary'].sum()

# Count: Number of employees earning more than 50000
high_earners_count = df[df['salary'] > 50000].shape[0]

# Limit: Get the top 5 highest earners
top_5_earners = df.nlargest(5, 'salary')

# Skip: Skip the first 2 employees
skipped_df = df.iloc[2:]

# Remove: Remove employees from a specific department
filtered_df = df[df['department'] != 'IT']

# Save the filtered result back to HDFS
filtered_json = filtered_df.to_json(orient='records')

try:
    with hdfs_client.write('/home/hadoop/filtered_employees.json', encoding='utf-8',
        overwrite=True) as writer:
        writer.write(filtered_json)
    print("Filtered JSON file saved successfully.")
except Exception as e:
    print(f'Error saving filtered JSON data: {e}')
exit(1)
```

```
# Print results

print(f'Projection: Select only name and salary columns')
print(f'{projected_df}')

print(f'Aggregation: Calculate total salary')

print(f'Total Salary: {total_salary}')
print(f'\n')

print(f'# Count: Number of employees earning more than 50000')

print(f'Number of High Earners (>50000):
{high_earners_count}') print(f'\n') print(f'limit Top 5 highest
salary')

print(f'Top 5 Earners: \n{top_5_earners}')
print(f'\n')
print(f'Skipped DataFrame (First 2 rows skipped): \n{skipped_df}')
print(f'\n')
print(f'Filtered DataFrame (Sales department removed): \n{filtered_df}')

run the file by bash: python3
process_data.py
```

output

```
hadoop@ubuntu:~/Ex6$ hdfs dfs -chmod 777 /ex6
hadoop@ubuntu:~/Ex6$ python3 process_data.py
Raw JSON Data: [
{"name": "John Doe", "age": 30, "department": "HR", "salary": 50000},
{"name": "Jane Smith", "age": 25, "department": "IT", "salary": 60000},
{"name": "Alice Johnson", "age": 35, "department": "Finance", "salary": 70000},
{"name": "Bob Brown", "age": 28, "department": "Marketing", "salary": 55000},
{"name": "Charlie Black", "age": 45, "department": "IT", "salary": 80000}
]
```

Filtered JSON file saved successfully.

Projection: Select only name and salary columns

| | name | salary |
|---|---------------|--------|
| 0 | John Doe | 50000 |
| 1 | Jane Smith | 60000 |
| 2 | Alice Johnson | 70000 |
| 3 | Bob Brown | 55000 |
| 4 | Charlie Black | 80000 |

Aggregation: Calculate total salary

Total Salary: 315000

Count: Number of employees earning more than 50000

Number of High Earners (>50000): 4

Limit: Top 5 highest salary

Top 5 Earners:

| | name | age | department | salary |
|---|---------------|-----|------------|--------|
| 4 | Charlie Black | 45 | IT | 80000 |
| 2 | Alice Johnson | 35 | Finance | 70000 |
| 1 | Jane Smith | 25 | IT | 60000 |
| 3 | Bob Brown | 28 | Marketing | 55000 |
| 0 | John Doe | 30 | HR | 50000 |

Limit: Top 5 highest salary

Top 5 Earners:

| | name | age | department | salary |
|---|---------------|-----|------------|--------|
| 4 | Charlie Black | 45 | IT | 80000 |
| 2 | Alice Johnson | 35 | Finance | 70000 |
| 1 | Jane Smith | 25 | IT | 60000 |
| 3 | Bob Brown | 28 | Marketing | 55000 |
| 0 | John Doe | 30 | HR | 50000 |

Skipped DataFrame (First 2 rows skipped):

| | name | age | department | salary |
|---|---------------|-----|------------|--------|
| 2 | Alice Johnson | 35 | Finance | 70000 |
| 3 | Bob Brown | 28 | Marketing | 55000 |
| 4 | Charlie Black | 45 | IT | 80000 |

Filtered DataFrame (IT department removed):

| | name | age | department | salary |
|---|---------------|-----|------------|--------|
| 0 | John Doe | 30 | HR | 50000 |
| 2 | Alice Johnson | 35 | Finance | 70000 |
| 3 | Bob Brown | 28 | Marketing | 55000 |

```
hadoop@ubuntu:~/Ex6$
```